



Regular Expressions in Java





Regular Expressions

- A **regular expression** is a kind of pattern that can be applied to text (**Strings**, in Java)
- A regular expression either **matches** the text (or part of the text), or it fails to match
 - If a regular expression matches a part of the text, then you can easily find out which part
 - If a regular expression is complex, then you can easily find out which parts of the regular expression match which parts of the text
 - With this information, you can readily extract parts of the text, or do substitutions in the text
- Regular expressions are an extremely useful tool for manipulating text
 - Regular expressions are heavily used in the automatic generation of Web pages



Perl and Java

- The **Perl** programming language is heavily used in server-side programming, because
 - Much server-side programming is text manipulation
 - Regular expressions are built into the syntax of Perl
- Beginning with Java 1.4, Java has a regular expression package, **java.util.regex**
 - Java's regular expressions are almost identical to those of Perl
 - This new capability greatly enhances Java 1.4's text handling
- Regular expressions in Java 1.4 are just a normal package, with no new syntax to support them
 - Java's regular expressions are just as powerful as Perl's, but
 - Regular expressions are easier and more convenient in Perl



A first example

- The regular expression "[a-z]+" will match a sequence of one or more lowercase letters
 - [a-z] means any character from a through z, inclusive
 - + means "one or more"
- Suppose we apply this pattern to the String "Now is the time"
 - There are *three ways* we can apply this pattern:
 - To the *entire string*: it fails to match because the string contains characters other than lowercase letters
 - To the *beginning of the string*: it fails to match because the string does not begin with a lowercase letter
 - To *search the string*: it will succeed and match **ow**
 - If applied repeatedly, it will find **is**, then **the**, then **time**, then fail



Doing it in Perl and Ruby

- In both Perl and Ruby, a regular expression is written between forward slashes, for example, `/[a-z]+/`
- Regular expressions are values, and can be used as such
 - For example, `line.split(/\s+/)`
- We can search for matches to a regular expression with the `=~` operator
 - For example, `name = "Dave"; name =~ /[a-z]/;` will find `ave`



Doing it in Java, I

- First, you must *compile* the pattern

```
import java.util.regex.*;  
Pattern p = Pattern.compile("[a-z]+");
```

- Next, you must create a *matcher* for a specific piece of text by sending a message to your pattern

```
Matcher m = p.matcher("Now is the time");
```

- Points to notice:

- **Pattern** and **Matcher** are both in **java.util.regex**
- Neither **Pattern** nor **Matcher** has a public constructor; you create these by using methods in the **Pattern** class
- The matcher contains information about *both* the pattern to use *and* the text to which it will be applied



Doing it in Java, II

- Now that we have a matcher `m`,
 - `m.matches()` returns `true` if the pattern matches the entire text string, and `false` otherwise
 - `m.lookingAt()` returns `true` if the pattern matches at the beginning of the text string, and `false` otherwise
 - `m.find()` returns `true` if the pattern matches any part of the text string, and `false` otherwise
 - If called again, `m.find()` will start searching from where the last match was found
 - `m.find()` will return `true` for as many matches as there are in the string; after that, it will return `false`
 - When `m.find()` returns `false`, matcher `m` will be *reset* to the beginning of the text string (and may be used again)



Finding what was matched

- *After a successful match, `m.start()` will return the index of the first character matched*
- *After a successful match, `m.end()` will return the index of the last character matched, *plus one**
- If no match was attempted, or if the match was unsuccessful, `m.start()` and `m.end()` will throw an `IllegalStateException`
 - This is a `RuntimeException`, so you don't have to catch it
- It may seem strange that `m.end()` returns the index of the last character matched plus one, but this is just what most String methods require
 - For example, `"Now is the time".substring(m.start(), m.end())` will return exactly the matched substring



A complete example

```
import java.util.regex.*;

public class RegexTest {
    public static void main(String args[]) {
        String pattern = "[a-z]+";
        String text = "Now is the time";
        Pattern p = Pattern.compile(pattern);
        Matcher m = p.matcher(text);
        while (m.find()) {
            System.out.print(text.substring(m.start(), m.end()) + "*");
        }
    }
}
```

Output: ow*is*the*time*



Additional methods

- If `m` is a matcher, then
 - `m.replaceFirst(replacement)` returns a new String where the first substring matched by the pattern has been replaced by *replacement*
 - `m.replaceAll(replacement)` returns a new String where every substring matched by the pattern has been replaced by *replacement*
 - `m.find(startIndex)` looks for the next pattern match, starting at the specified index
 - `m.reset()` resets this matcher
 - `m.reset(newText)` resets this matcher and gives it new text to examine (which may be a `String`, `StringBuffer`, or `CharBuffer`)



Some simple patterns

`abc` exactly this sequence of three letters

`[abc]` any *one* of the letters `a`, `b`, or `c`

`[^abc]` any character *except* one of the letters `a`, `b`, or `c`
(immediately within an open bracket, `^` means “not,”
but anywhere else it just means the character `^`)

`[a-z]` any *one* character from `a` through `z`, inclusive

`[a-zA-Z0-9]` any *one* letter or digit



Sequences and alternatives

- If one pattern is followed by another, the two patterns must match consecutively
 - For example, `[A-Za-z]+[0-9]` will match one or more letters immediately followed by one digit
- The vertical bar, `|`, is used to separate alternatives
 - For example, the pattern `abc|xyz` will match either `abc` or `xyz`



Some predefined character classes

`.` any one character except a line terminator

`\d` a digit: `[0-9]`

`\D` a non-digit: `[^0-9]`

`\s` a whitespace character: `[\t\n\x0B\f\r]`

`\S` a non-whitespace character: `[^\s]`

`\w` a word character: `[a-zA-Z_0-9]`

`\W` a non-word character: `[^\w]`

Notice the space.
Spaces are **significant**
in regular expressions!



Boundary matchers

- These patterns match the *empty string* if at the specified position:
 - ^ the beginning of a line
 - \$ the end of a line
 - \b a word boundary
 - \B not a word boundary
 - \A the beginning of the input (can be multiple lines)
 - \Z the end of the input except for the final terminator, if any
 - \z the end of the input



Greedy quantifiers

(The term “greedy” will be explained later)

Assume X represents some pattern

$X?$ optional, X occurs once or not at all

X^* X occurs zero or more times

X^+ X occurs one or more times

$X\{n\}$ X occurs exactly n times

$X\{n,\}$ X occurs n or more times

$X\{n,m\}$ X occurs at least n but not more than m times

Note that these are all *postfix* operators, that is, they come *after* the operand



Types of quantifiers

- A **greedy quantifier** will match as much as it can, and back off if it needs to
 - We'll do examples in a moment
- A **reluctant quantifier** will match as little as possible, then take more if it needs to
 - You make a quantifier reluctant by appending a **?**:
 $X??$ $X^*?$ $X+?$ $X\{n\}?$ $X\{n,\}?$ $X\{n,m\}?$
- A **possessive quantifier** will match as much as it can, and never let go
 - You make a quantifier possessive by appending a **+**:
 $X?+$ X^*+ $X++$ $X\{n\}+$ $X\{n,\}+$ $X\{n,m\}+$



Quantifier examples

- Suppose your text is **aardvark**
 - Using the pattern **a*ardvark** (**a*** is greedy):
 - The **a*** will first match **aa**, but then **ardvark** won't match
 - The **a*** then “backs off” and matches only a single **a**, allowing the rest of the pattern (**ardvark**) to succeed
 - Using the pattern **a*?ardvark** (**a*?** is reluctant):
 - The **a*?** will first match zero characters (the null string), but then **ardvark** won't match
 - The **a*?** then extends and matches the first **a**, allowing the rest of the pattern (**ardvark**) to succeed
 - Using the pattern **a*+ardvark** (**a*+** is possessive):
 - The **a*+** will match the **aa**, and will not back off, so **ardvark** never matches and the pattern match fails



Capturing groups

- In regular expressions, parentheses are used for grouping, but they also **capture** (keep for later use) anything matched by that part of the pattern
 - Example: `([a-zA-Z]*)([0-9]*)` matches any number of letters followed by any number of digits
 - If the match succeeds, `\1` holds the matched letters and `\2` holds the matched digits
 - In addition, `\0` holds everything matched by the entire pattern
- Capturing groups are numbered by counting their *opening parentheses* from left to right:
 - `((A)(B(C)))`
`1 2 3 4`
`\0 = \1 = ((A)(B(C))), \2 = (A), \3 = (B(C)), \4 = (C)`
- Example: `([a-zA-Z])\1` will match a double letter, such as letter



Capturing groups in Java

- If `m` is a matcher that has just performed a successful match, then
 - `m.group(n)` returns the String matched by capturing group *n*
 - This could be an empty string
 - This will be `null` if the pattern as a whole matched but this particular group didn't match anything
 - `m.group()` returns the String matched by the entire pattern (same as `m.group(0)`)
 - This could be an empty string
- If `m` didn't match (or wasn't tried), then these methods will throw an `IllegalStateException`



Example use of capturing groups

- Suppose **word** holds a word in English
- Also suppose we want to move all the consonants at the beginning of **word** (if any) to the end of the word (so **string** becomes **ingstr**)
 - ```
Pattern p = Pattern.compile("([^aeiou]*)(.*)");
Matcher m = p.matcher(word);
if (m.matches()) {
 System.out.println(m.group(2) + m.group(1));
}
```
- Note the use of **(.\*)** to indicate “all the rest of the characters”



# Double backslashes

---

- Backslashes have a special meaning in regular expressions; for example, `\b` means a word boundary
- Backslashes have a special meaning in Java; for example, `\b` means the backspace character
- Java syntax rules apply first!
  - If you write `"\b[a-z]+\b"` you get a string with backspace characters in it--this is *not* what you want!
  - Remember, you can quote a backslash with another backslash, so `"\\b[a-z]+\\b"` gives the correct string
- Note: if you *read in* a String from somewhere, this does not apply--you get whatever characters are actually there



# Escaping metacharacters

---

- A lot of special characters--parentheses, brackets, braces, stars, plus signs, etc.--are used in defining regular expressions; these are called **metacharacters**
- Suppose you want to search for the character sequence **a**\* (an **a** followed by a star)
  - "a\*"; doesn't work; that means "zero or more **a**s"
  - "a\\*"; doesn't work; since a star doesn't *need* to be escaped (in Java String constants), Java just ignores the \
  - "a\\\*" *does* work; it's the three-character string **a**, \, \*
- Just to make things even more difficult, it's *illegal* to escape a *non*-metacharacter in a regular expression



# Spaces

---

- There is only one thing to be said about spaces (blanks) in regular expressions, but it's important:
  - *Spaces are significant!*
- A space stands for a *space*--when you put a space in a pattern, that means to match a space in the text string
- It's a *really bad idea* to put spaces in a regular expression just to make it look better



# Additions to the String class

---

- All of the following are **public**:
  - `public boolean matches(String regex)`
  - `public String replaceFirst(String regex, String replacement)`
  - `public String replaceAll(String regex, String replacement)`
  - `public String[] split(String regex)`
  - `public String[] split(String regex, int limit)`
    - If the limit *n* is greater than zero then the pattern will be applied at most *n* - 1 times, the array's length will be no greater than *n*, and the array's last entry will contain all input beyond the last matched delimiter.
    - If *n* is non-positive then the pattern will be applied as many times as possible





# Thinking in regular expressions

---

- Regular expressions are *not* easy to use at first
  - It's a bunch of punctuation, not words
  - The individual pieces are not hard, but it takes practice to learn to put them together correctly
  - Regular expressions form a miniature programming language
    - It's a different kind of programming language than Java, and requires you to learn new thought patterns
  - In Java you can't just *use* a regular expression; you have to first create Patterns and Matchers
  - Java's syntax for String constants doesn't help, either
- Despite all this, regular expressions bring so much power and convenience to String manipulation that they are well worth the effort of learning



# The End

---

“A little learning is a dangerous thing;  
drink deep, or taste not the Pierian spring:  
there shallow draughts intoxicate the brain,  
and drinking largely sobers us again.”

--Alexander Pope